



Content  $\oplus$  identity: each token's layer-4 residual split into a per-compartment constant (identity) and the rest (content)  
top: one sheet per compartment, a thread per token running through the stack -- vertical threads mean the pages are congruent  
bottom: where those sheets actually sit along the identity axis, on one shared scale

Default init  
(no alignment pressure)  
cos 0.132  $\rightarrow$  0.057 with identity removed



Init-copy  
(embeddings + head shared at init)  
cos 0.981  $\rightarrow$  0.996 with identity removed



InfoNCE  
(explicit layer-4 alignment)  
cos 0.963  $\rightarrow$  0.984 with identity removed

